

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

"JnanaSangama", Belgaum -590014, Karnataka.



LAB REPORT on

MACHINE LEARNING

Submitted by

HARSHITHA RM (1BM19CS060)

in partial fulfillment for the award of the degree of
BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING
BENGALURU-560019 May-2022 to July-2022
(Autonomous Institution under VTU)

B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “MACHINE LEARNING” carried out by **HARSHITHA RM (1BM19CS060)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2022. The Lab report has been approved as it satisfies the academic requirements in respect of a **Machine Learning - (20CS6PCMAL)** work prescribed for the said degree.

Name of the Lab-Incharge:

Dr G R ASHA

Assistant Professor
Department of CSE
BMSCE, Bengaluru

Dr. Jyothi S Nayak

Professor and Head
Department of CSE
BMSCE, Bengaluru

Index Sheet

Sl. No.	Experiment Title	Page No.
1	FIND-S Algorithm	4
2	Candidate Elimination Algorithm	6
3	Decision tree based ID3 Algorithm	9
4	Linear Regression	17
5	Naive Bayesian Classifier	20
6	Bayesian Network	23
7	KMeans Clustering Algorithm	26
8	EM Algorithm	29
9	KNN Algorithm	31
10	Weighted Linear Regression Algorithm	34

Course Outcome

--	--

EXPERIMENT 1

Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples.

CODE:

```
import csv
import pandas as pd
import numpy as np

data = pd.read_csv("Desktop/data.csv")
print(data, "\n")

#array of all the attributes d =
np.array(data)[:,-1]
print("\n
The attributes are: ",d)

target = np.array(data)[:,-1]
print("\n
The target is: ",target)

def findS(c,t):
    for i, val in enumerate(t):
        if val == "Yes":
            specific_hypothesis = c[i].copy()
            break

    for i, val in enumerate(c):
        if t[i] == "Yes":
            for x in range(len(specific_hypothesis)):
                if val[x] != specific_hypothesis[x]:
                    specific_hypothesis[x] = '?'
            else:
                pass
```

```
return specific_hypothesis
```

```
print("\n The final hypothesis is:",findS(d,target))
```

OUTPUT:

	Weather	Temperature	Humidity	Wind	Goes
0	Sunny	Warm	Mild	Strong	Yes
1	Rainy	Cold	Mild	Normal	No
2	Sunny	Moderate	Normal	Normal	Yes
3	Sunny	Cold	High	Strong	Yes

```
The attributes are: [['Sunny' 'Warm' 'Mild' 'Strong']  
['Rainy' 'Cold' 'Mild' 'Normal']  
['Sunny' 'Moderate' 'Normal' 'Normal']  
['Sunny' 'Cold' 'High' 'Strong']]
```

```
The target is: ['Yes' 'No' 'Yes' 'Yes']
```

```
The final hypothesis is: ['Sunny' '?' '?' '?']
```

EXPERIMENT 2

For a given set of training data examples stored in a .CSV file, implement and demonstrate the Candidate-Elimination algorithm to output a description of the set of all hypotheses consistent with the training examples.

CODE:

```
import numpy as np  
import pandas as pd
```

```

data = pd.read_csv('Desktop/shape.csv')
concepts = np.array(data.iloc[:,0:-1])
print("\nInstances are:\n",concepts) target
= np.array(data.iloc[:, -1]) print("\nTarget
Values are: ",target)

def learn(concepts, target):
    specific_h = concepts[0].copy() print("\nInitialization of specific_h and
general_h") print("\nSpecific Boundary: ", specific_h) general_h = ["?"
for i in range(len(specific_h))] for i in range(len(specific_h)) print("\nGeneric
Boundary: ",general_h)

    for i, h in enumerate(concepts):
        print("\nInstance", i+1 , "is ", h) if
        target[i] == "yes":
            print("Instance is Positive ")
            for x in range(len(specific_h)):
                if h[x]!= specific_h[x]:
                    specific_h[x] ='?'
                    general_h[x][x] ='?'

            if target[i] == "no":
                print("Instance is Negative ")
                for x in range(len(specific_h)):
                    if h[x]!= specific_h[x]:
                        general_h[x][x] = specific_h[x]
                    else:
                        general_h[x][x] = '?'

```

```

        print("Specific Boundary after ", i+1, "Instance is ", specific_h)
    print("Generic Boundary after ", i+1, "Instance is ", general_h)

    print("\n")

    indices = [i for i, val in enumerate(general_h) if val == ['?', '?', '?', '?', '?', '?']]
    for i in indices:
        general_h.remove(['?', '?', '?', '?', '?', '?'])

    return specific_h, general_h

s_final, g_final = learn(concepts, target)

print("Final Specific_h: ", s_final, sep="\n") print("Final
General_h: ", g_final, sep="\n")

```

OUTPUT:

```

In [3]: data = pd.read_csv('Desktop/shape.csv')
        concepts = np.array(data.iloc[:,0:-1])
        print("\nInstances are:\n",concepts)
        target = np.array(data.iloc[:,-1])
        print("\nTarget Values are: ",target)

```

Instances are:

```

[['big' 'red' 'circle']
 ['small' 'red' 'triangle']
 ['small' 'red' 'circle']
 ['big' 'blue' 'circle']
 ['small' 'blue' 'circle']]

```

Target Values are: ['no' 'no' 'yes' 'no' 'yes']

```

Initialization of specific_h and general_h

Specific Boundary: ['big' 'red' 'circle']

Generic Boundary: [['?', '?', '?'], ['?', '?', '?'], ['?', '?', '?']]

Instance 1 is ['big' 'red' 'circle']
Instance is Negative
Specific Boundary after 1 Instance is ['big' 'red' 'circle']
Generic Boundary after 1 Instance is [['?', '?', '?'], ['?', '?', '?'], ['?', '?', '?']]

Instance 2 is ['small' 'red' 'triangle']
Instance is Negative
Specific Boundary after 2 Instance is ['big' 'red' 'circle']
Generic Boundary after 2 Instance is [['big', '?', '?'], ['?', '?', '?'], ['?', '?', 'circle']]

Instance 3 is ['small' 'red' 'circle']
Instance is Positive
Specific Boundary after 3 Instance is ['?' 'red' 'circle']
Generic Boundary after 3 Instance is [['?', '?', '?'], ['?', '?', '?'], ['?', '?', 'circle']]

Instance 4 is ['big' 'blue' 'circle']
Instance is Negative
Specific Boundary after 4 Instance is ['?' 'red' 'circle']
Generic Boundary after 4 Instance is [['?', '?', '?'], ['?', 'red', '?'], ['?', '?', '?']]

Instance 5 is ['small' 'blue' 'circle']
Instance is Positive
Specific Boundary after 5 Instance is ['?' '?' 'circle']
Generic Boundary after 5 Instance is [['?', '?', '?'], ['?', '?', '?'], ['?', '?', '?']]

Final Specific_h:
['?' '?' 'circle']
Final General_h:
[['?', '?', '?'], ['?', '?', '?'], ['?', '?', '?']]

```

EXPERIMENT 3

Write a program to demonstrate the working of the decision tree based ID3 algorithm. Use an appropriate data set for building the decision tree and apply this knowledge to classify a new sample.

CODE:

WITHOUT ALGO:

```

import csv
def load_csv(filename):
    lines=csv.reader(open(filename,"r"));

    dataset = list(lines)

```



```
headers = dataset.pop(0)
return dataset,headers
```

```
class Node:    def
__init__(self,attribute):
self.attribute=attribute
self.children=[]
self.answer=""
```

```
def subtables(data,col,delete):
    dic={}
    coldata=[row[col] for row in data]
    attr=list(set(coldata))
```

```
    counts=[0]*len(attr)
    r=len(data)
    c=len(data[0])    for x in
range(len(attr)):    for y
in range(r):        if
data[y][col]==attr[x]:
        counts[x]+=1
```

```
    for x in range(len(attr)):
        dic[attr[x]]=[[0 for i in range(c)] for j in range(counts[x])]
    pos=0    for y in range(r):
        if data[y][col]==attr[x]:
    if delete:
        del data[y][col]
    dic[attr[x]][pos]=data[y]        pos+=1
```

```
    return attr,dic
```

```
def entropy(S):
```

```
    attr=list(set(S))    if
```

```
    len(attr)==1:
```

```
    return 0
```

```
    counts=[0,0]
```

```
    for i in range(2):
```

```
        counts[i]=sum([1 for x in S if attr[i]==x])/(len(S)*1.0)
```

```
    sums=0    for cnt in counts:
```

```
        sums+=-1*cnt*math.log(cnt,2)
```

```
    return sums
```

```
def compute_gain(data,col):
```

```
    attr,dic = subtables(data,col,delete=False)
```

```
    total_size=len(data)
```

```
    entropies=[0]*len(attr)    ratio=[0]*len(attr)
```

```
    total_entropy=entropy([row[-1] for row in data])
```

```
    for x in range(len(attr)):
```

```
        ratio[x]=len(dic[attr[x]])/(total_size*1.0)
```

```
    entropies[x]=entropy([row[-1] for row in dic[attr[x]]])
```

```
    total_entropy-=ratio[x]*entropies[x]    return
```

```
    total_entropy
```

```

def build_tree(data,features):
    lastcol=[row[-1] for row in data]
    if(len(set(lastcol))==1:
        node=Node("")
        node.answer=lastcol[0]    return
    node

    n=len(data[0])-1
    gains=[0]*n    for
    col in range(n):
        gains[col]=compute_gain(data,col)
    split=gains.index(max(gains))
    node=Node(features[split])    fea =
    features[:split]+features[split+1:]
    attr,dic=subtables(data,split,delete=True
    )

    for x in range(len(attr)):
        child=build_tree(dic[attr[x]],fea)
    node.children.append((attr[x],child))    return
    node

def print_tree(node,level):
    if node.answer!="":
        print(" "*level,node.answer)
        return

    print(" "*level,node.attribute)
    for value,n in node.children:

```

```
print("        "*(level+1),value)
print_tree(n,level+2)
```

```
def classify(node,x_test,features):
    if node.answer!="":
        print(node.answer)
        return
    pos=features.index(node.attribute)
    for value, n in node.children:    if
    x_test[pos]==value:
        classify(n,x_test,features)
```

```
'''Main program''' dataset,features=load_csv("data.csv") node1=build_tree(dataset,features)
```

```
print("The decision tree for the dataset using ID3 algorithm is")
print_tree(node1,0) testdata,features=load_csv("test.csv")
```

```
for xtest in testdata:
    print("The test instance:",xtest)
    print("The label for test instance:",end=" ")
    classify(node1,xtest,features)
```

WITH ALGO:

```
import numpy as np
import pandas as pd
import math
```

```
data = pd.DataFrame(data=pd.read_csv('data.csv')) print(data)
```

```

def countPosNeg(data):
    pos = data.iloc[:, -1:].value_counts()['yes']
    neg = len(data) - pos    return pos, neg

def calcEntropy(pos, neg):
    entropy = -(pos/(pos+neg))*math.log2(pos/(pos+neg)) -(neg/(pos+neg))*math.log2(neg/(pos+neg))

    return entropy

def calcAverageInformation(data):    #
    iterate through each attribute (col)
    attribs = data.iloc[:, :-1].columns.values
    print(attribs)

    for attrib in attribs:    # get
    possible values    values =
    data[attrib].unique()

    valueEntropies = pd.DataFrame(0, columns=['p', 'n', 'entropy'], index=values)
    print()
    print(attrib)
    print(valueEntropies)

    # iterate through whole dataframe
    for i in data.index:
        print(data['Answer'][i])
    if data['Answer'][i] == 'yes':

```

```

        valueEntropies[data[attrib]]['p'] += 1
elif data['Answer'][i] == 'no':
    valueEntropies[data[attrib]]['n'] += 1

for value in valueEntropies:
    value['entropy'] = calcEntropy(value['p'], value['n'])

print(valueEntropies)

return 10

calcAverageInformation(data)

def calcGain(entropy, avg_info):
    return entropy - avg_info

# data for the total dataset

tot_pos, tot_neg = countPosNeg(data) tot_entropy
= calcEntropy(tot_pos, tot_neg) print(tot_entropy)

# iterate through dataset and calc pos, neg and entropy vals for each column

```

OUTPUT:

The decision tree for the dataset using ID3 algorithm is

```
Outlook
  sunny
    Humidity
      normal
        yes
      high
        no
  rain
    Wind
      weak
        yes
      strong
        no
  overcast
    yes
```

```
import numpy as np
import pandas as pd
import math

data = pd.DataFrame(data=pd.read_csv('Desktop/data.csv'))
print(data)

# print(data['Answer'])
```

	Outlook	Temperature	Humidity	Wind	Answer
0	sunny	hot	high	weak	no
1	sunny	hot	high	strong	no
2	overcast	hot	high	weak	yes
3	rain	mild	high	weak	yes
4	rain	cool	normal	weak	yes
5	rain	cool	normal	strong	no
6	overcast	cool	normal	strong	yes
7	sunny	mild	high	weak	no
8	sunny	cool	normal	weak	yes
9	rain	mild	normal	weak	yes
10	sunny	mild	normal	strong	yes
11	overcast	mild	high	strong	yes
12	overcast	hot	normal	weak	yes
13	rain	mild	high	strong	no

```

        valueEntropies[data[attrib]][p] += 1
    elif data['Answer'][i] == 'no':
        valueEntropies[data[attrib]][n] += 1

    for value in valueEntropies:
        value['entropy'] = calcEntropy(value[p], value[n])

    print(valueEntropies)

    # print(data['Outlook'].unique())

    return 10

calcAverageInformation(data)

['Outlook' 'Temperature' 'Humidity' 'Wind']

Outlook
      p  n  entropy
sunny  0  0        0
overcast 0  0        0
rain    0  0        0
no

def calcGain(entropy, avg_info):
    return entropy - avg_info

# data for the total dataset
tot_pos, tot_neg = countPosNeg(data)
tot_entropy = calcEntropy(tot_pos, tot_neg)
print(tot_entropy)

# iterate through dataset and calc pos, neg and entropy vals for each attribute

Answer
yes    0.940286
dtype: float64

```

EXPERIMENT 4

Implement the Linear Regression algorithm in order to fit data points. Select appropriate data set for your experiment and draw graphs.

CODE:

```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

```



```
dataset = pd.read_csv('salary_data.csv')
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, 1].values

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/3, random_state=0)

# Fitting Simple Linear Regression to the Training set
from sklearn.linear_model import LinearRegression
regressor = LinearRegression() regressor.fit(X_train,
y_train)

# Predicting the Test set results y_pred
= regressor.predict(X_test)

# Visualizing the Training set results viz_train
= plt viz_train.scatter(X_train, y_train,
color='red') viz_train.plot(X_train,
regressor.predict(X_train), color='blue')
viz_train.title('Salary VS Experience (Training
set)') viz_train.xlabel('Year of Experience')
viz_train.ylabel('Salary') viz_train.show()

# Visualizing the Test set results viz_test = plt
viz_test.scatter(X_test, y_test, color='red')
viz_test.plot(X_train, regressor.predict(X_train), color='blue')
viz_test.title('Salary VS Experience (Test set)')
viz_test.xlabel('Year of Experience') viz_test.ylabel('Salary')
viz_test.show()
```

OUTPUT:

```
regressor.fit(X_train, y_train)
```

```
Out[5]: LinearRegression()
```

```
In [9]: # Predicting the Test set results
y_pred = regressor.predict(X_test)
print(y_pred)
```

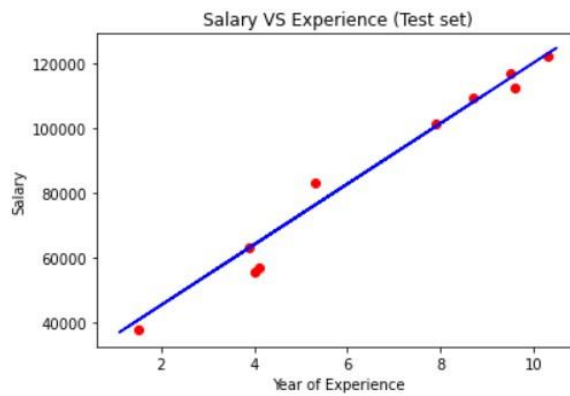
```
[ 40835.10590871 123079.39940819  65134.55626083  63265.36777221
 115602.64545369 108125.8914992  116537.23969801  64199.96201652
 76349.68719258 100649.1375447 ]
```

```
In [7]: # Visualizing the Training set results
viz_train = plt
viz_train.scatter(X_train, y_train, color='red')
viz_train.plot(X_train, regressor.predict(X_train), color='blue')
viz_train.title('Salary VS Experience (Training set)')
viz_train.xlabel('Year of Experience')
viz_train.ylabel('Salary')
viz_train.show()
```



Year of Experience

```
In [8]: # Visualizing the Test set results
viz_test = plt
viz_test.scatter(X_test, y_test, color='red')
viz_test.plot(X_train, regressor.predict(X_train), color='blue')
viz_test.title('Salary VS Experience (Test set)')
viz_test.xlabel('Year of Experience')
viz_test.ylabel('Salary')
viz_test.show()
```



EXPERIMENT 5

Write a program to implement the naïve Bayesian classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering few test data sets

CODE:

```
import pandas as pd from sklearn.model_selection
import train_test_split from sklearn.naive_bayes
import GaussianNB from sklearn import metrics

df = pd.read_csv("Downloads/data.csv")

feature_col_names = ['num_preg', 'glucose_conc', 'diastolic_bp', 'thickness', 'insulin', 'bmi', 'diab_pred',
'age']

predicted_class_names = ['diabetes']

X = df[feature_col_names].values y =
df[predicted_class_names].values
```

```
print(df.head) xtrain,xtest,ytrain,ytest=train_test_split(X,y,test_size=0.40)
```

```
print ('\n the total number of Training Data :',ytrain.shape) print  
('\n the total number of Test Data :',ytest.shape)
```

```
clf = GaussianNB().fit(xtrain,ytrain.ravel()) predicted  
= clf.predict(xtest) predictTestData=  
clf.predict([[6,148,72,35,0,33.6,0.627,50]])
```

```
print('\n Confusion matrix') print(metrics.confusion_matrix(ytest,predicted))
```

```
print('\n Accuracy of the classifier is',metrics.accuracy_score(ytest,predicted))
```

```
print('\n The value of Precision', metrics.precision_score(ytest,predicted))
```

```
print('\n The value of Recall', metrics.recall_score(ytest,predicted))
```

```
print("Predicted Value for individual Test Data:", predictTestData)
```

OUTPUT:

	<bound	method	NDFrame.head of	num_preg	glucose_conc	diastolic_bp	thickness	insulin	bmi \
0		6	148	72	35	0	33.6		
1		1	85	66	29	0	26.6		
2		8	183	64	0	0	23.3		
3		1	89	66	23	94	28.1		
4		0	137	40	35	168	43.1		
...			
140		3	128	78	0	0	21.1		
141		5	106	82	30	0	39.5		
142		2	108	52	26	63	32.5		
143		10	108	66	0	0	32.4		
144		4	154	62	31	284	32.8		

	diab_pred	age	diabetes
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
...
140	0.268	55	0
141	0.286	38	0
142	0.318	22	0
143	0.272	42	1
144	0.237	23	0

[145 rows x 9 columns]>

```
print('\n the total number of Training Data :',ytrain.shape)
print ('\n the total number of Test Data :',ytest.shape)
```

the total number of Training Data : (87, 1)

the total number of Test Data : (58, 1)

```
predictTestData= clf.predict([[6,148,72,35,0,33.6,0.627,50]])
```

```
In [7]: ► print('\n Confusion matrix')
print(metrics.confusion_matrix(ytest,predicted))
```

```
Confusion matrix
[[32  6]
 [12  8]]
```

```
In [8]: ► print('\n Accuracy of the classifier is',metrics.accuracy_score(ytest,predicted))
print('\n The value of Precision', metrics.precision_score(ytest,predicted))
print('\n The value of Recall', metrics.recall_score(ytest,predicted))
print("Predicted Value for individual Test Data:", predictTestData)
```

Accuracy of the classifier is 0.6896551724137931

The value of Precision 0.5714285714285714

The value of Recall 0.4

Predicted Value for individual Test Data: [1]

Experiment 6

Write a program to construct a Bayesian network considering training data. Use this model to make predictions.

CODE:

```
import numpy as np
import pandas as pd
import csv

import pgmpy.estimators
import pgmpy.models
import pgmpy.inference

MaximumLikelihoodEstimator = pgmpy.estimators.MaximumLikelihoodEstimator
BayesianModel = pgmpy.models.BayesianModel
VariableElimination = pgmpy.inference.VariableElimination
```

```

#read Cleveland Heart Disease data heartDisease
= pd.read_csv('Downloads/data.csv') heartDisease
= heartDisease.replace('?',np.nan)

#display the data print('Sample instances from the dataset
are given below') print(heartDisease.head())

#display the Attributes names and datatypes print('\n
Attributes and datatypes') print(heartDisease.dtypes)

#Create Model-Bayesian Network
model =
BayesianModel([('age','heartDisease'),('sex','heartDisease'),('exang','heartDisease'),('cp','heartDisease'),('
'restecg','heartDisease'),('heartDisease','chol')])

#Learning CPDs using Maximum Likelihood Estimators print('\n
Learning CPD using Maximum likelihood estimators')
model.fit(heartDisease,estimator=MaximumLikelihoodEstimator)

#Inferencing with Bayesian Network print('\n
Inferencing with Bayesian Network:')
heartDiseasetest_infer = VariableElimination(model)

#computing the Probability of heartDisease given restecg print('\n 1.Probability of
heartDisease given evidence= restecg :1')
q1=heartDiseasetest_infer.query(variables=['heartDisease'],evidence={'restecg':1})
print(q1)

#computing the Probability of heartDisease given cp print('\n 2.Probability of
heartDisease given evidence= cp:2 ')

```

```
q2=heartDiseasetest_infer.query(variables=['heartDisease'],evidence={'cp':2})  
print(q2)
```

OUTPUT:

Sample instances from the dataset are given below

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	1	145	233	1	2	150	0	2.3	3	
1	67	1	4	160	286	0	2	108	1	1.5	2	
2	67	1	4	120	229	0	2	129	1	2.6	2	
3	37	1	3	130	250	0	0	187	0	3.5	3	
4	41	0	2	130	204	0	2	172	0	1.4	1	

	ca	thal	heartDisease
0	0	6	0
1	3	3	2
2	2	7	1
3	0	3	0
4	0	3	0

Attributes and datatypes

```
age          int64
sex          int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
exang        int64
oldpeak      float64
slope        int64
ca           int64
thal         int64
heartDisease int64
dtype: object
```

Learning CPD using Maximum likelihood estimators

Inferencing with Bayesian Network:

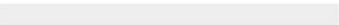
1.Probability of heartDisease given evidence= restecg :1

Finding Elimination Order: : 0%  0/4 [00:00<?, ?it/s]

Eliminating: cp: 100%  4/4 [00:00<00:00, 41.78it/s]

heartDisease	phi(heartDisease)
heartDisease(0)	0.1972
heartDisease(1)	0.1970
heartDisease(2)	0.1976
heartDisease(3)	0.1976
heartDisease(4)	0.2106

2.Probability of heartDisease given evidence= cp:2

Finding Elimination Order: : 0%  0/4 [00:00<?, ?it/s]

Eliminating: restecg: 100%  4/4 [00:00<00:00, 72.92it/s]

heartDisease	phi(heartDisease)
heartDisease(0)	0.3138
heartDisease(1)	0.2150
heartDisease(2)	0.1552
heartDisease(3)	0.1633
heartDisease(4)	0.1527

Experiment 7

Apply EM algorithm to cluster a set of data stored in a .CSV file. Compare the results of k-Means algorithm and EM algorithm.

CODE:

```
import matplotlib import
matplotlib.pyplot as plt import
seaborn as sns; sns.set() import
numpy as np

from sklearn.datasets import make_blobs X, y_true
= make_blobs(n_samples=300, centers=4,
              cluster_std=0.60, random_state=0) plt.scatter(X[:,
0], X[:, 1], s=50)

from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=4)
kmeans.fit(X) y_kmeans =
kmeans.predict(X)

plt.scatter(X[:, 0], X[:, 1], c=y_kmeans, s=50, cmap='viridis')

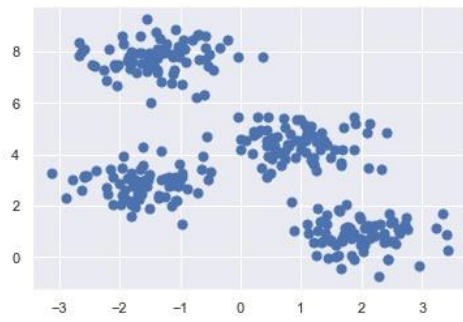
centers = kmeans.cluster_centers_
plt.scatter(centers[:, 0], centers[:, 1], c='black', s=200, alpha=0.5)

import pandas as pd import
numpy as np heartDisease
```

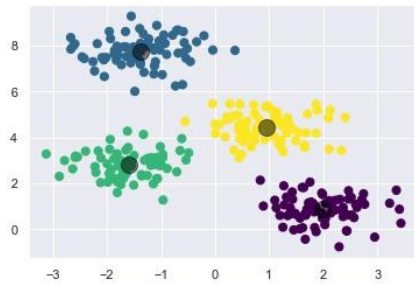
```
=  
pd.read_csv('Downloads/d  
ata.csv') heartDisease =  
heartDisease.replace('?',np.  
nan)  
  
heartDisease.head()  
  
trestbpsX = heartDisease.loc[:, 'trestbps'] cholY  
= heartDisease.loc[:, 'chol']  
plt.scatter(trestbpsX, cholY, s=50)  
  
kmeans2 = KMeans(n_clusters=2)  
combined_list = list(zip(trestbpsX, cholY))  
kmeans2.fit(combined_list) y_kmeans2 =  
kmeans2.predict(combined_list)  
  
plt.scatter(trestbpsX, cholY, c=y_kmeans2, s=50, cmap='viridis')  
  
centers = kmeans2.cluster_centers_  
plt.scatter(centers[:, 0], centers[:, 1], c='black', s=200, alpha=0.5)
```

OUTPUT:

Out[2]: <matplotlib.collections.PathCollection at 0x2006b964490>



Out[4]: <matplotlib.collections.PathCollection at 0x2006bc88610>

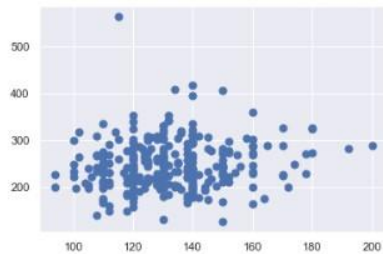


```
In [6]: import pandas as pd
import numpy as np
heartDisease = pd.read_csv('Downloads/data.csv')
heartDisease = heartDisease.replace('?', np.nan)
heartDisease.head()
```

Out[6]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	heartDisease
0	63	1	1	145	233	1	2	150	0	2.3	3	0	6	0
1	67	1	4	160	286	0	2	108	1	1.5	2	3	3	2
2	67	1	4	120	229	0	2	129	1	2.6	2	2	7	1
3	37	1	3	130	250	0	0	187	0	3.5	3	0	3	0
4	41	0	2	130	204	0	2	172	0	1.4	1	0	3	0

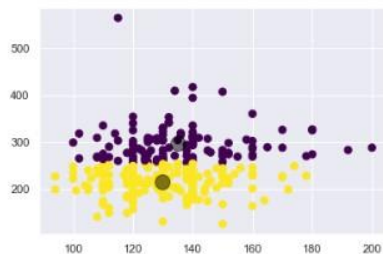
Out[8]: <matplotlib.collections.PathCollection at 0x2006c47ac40>



```
In [9]: kmeans2 = KMeans(n_clusters=2)
combined_list = list(zip(trestbpsX, cholY))
kmeans2.fit(combined_list)
y_kmeans2 = kmeans2.predict(combined_list)
```

```
In [10]: plt.scatter(trestbpsX, cholY, c=y_kmeans2, s=50, cmap='viridis')
centers = kmeans2.cluster_centers_
plt.scatter(centers[:, 0], centers[:, 1], c='black', s=200, alpha=0.5)
```

Out[10]: <matplotlib.collections.PathCollection at 0x2006c4d7d00>



Experiment 8

Write a program to implement k-Means algorithm to classify the iris data set. Print both correct and wrong predictions.

CODE:

```
from sklearn import datasets from
sklearn.cluster import KMeans from
sklearn.utils import shuffle import
numpy as np import pandas as pd

iris=datasets.load_iris()
X=iris.data
Y=iris.target

#Shuffle of Data
X,Y = shuffle(X,Y)

model=KMeans(n_clusters=3,init='k-means++',max_iter=10,n_init=1,random_state=3425)

#Training of the model model.fit(X)

# This is what KMeans thought (Prediction)
Y_Pred=model.labels_

from sklearn.metrics import confusion_matrix

cm=confusion_matrix(Y,Y_Pred) print(cm)

from sklearn.metrics import accuracy_score

print(accuracy_score(Y,Y_Pred))
```

```
#Defining EM Model from sklearn.mixture import
GaussianMixture
model2=GaussianMixture(n_components=3,random_state=3425)

#Training of the model model2.fit(X)

#Predicting classes for our data
Y_predict2= model2.predict(X)

#Accuracy of EM Model from sklearn.metrics
import confusion_matrix

cm=confusion_matrix(Y,Y_predict2) print(cm)

from sklearn.metrics import accuracy_score

print(accuracy_score(Y,Y_predict2))
```

OUTPUT:

```
[[ 0 50  0]
 [ 3  0 47]
 [36  0 14]]
0.09333333333333334
```

```
In [17]: ▶ #Defining EM Model
from sklearn.mixture import GaussianMixture
model2=GaussianMixture(n_components=3,random_state=3425)

#Training of the model
model2.fit(X)
```

```
Out[17]: GaussianMixture
GaussianMixture(n_components=3, random_state=3425)
```

```
In [18]: ▶ #Predicting classes for our data
Y_predict2= model2.predict(X)

#Accuracy of EM Model
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(Y,Y_predict2)
print(cm)

from sklearn.metrics import accuracy_score

print(accuracy_score(Y,Y_predict2))

[[ 0 50  0]
 [ 5  0 45]
 [50  0  0]]
0.0
```

Experiment 9

Write a program to implement k-Nearest Neighbor algorithm to classify the iris data set. Print both correct and wrong predictions.

CODE:


```
from sklearn.model_selection import train_test_split from
sklearn.neighbors import KNeighborsClassifier from sklearn.metrics
import classification_report, confusion_matrix from sklearn import
datasets
```

```
iris = datasets.load_iris()
```

```
X = iris.data
```

```
Y = iris.target
```

```
print('sepal-length','sepal-width','petal-length','petal-width')
```

```
print(X) print('target') print(Y)
```

```
x_train, x_test, y_train, y_test = train_test_split(X,Y,test_size=0.3)
```

```
classier = KNeighborsClassifier(n_neighbors=5)
```

```
classier.fit(x_train, y_train)
```

```
y_pred=classier.predict(x_test)
```

```
print('confusion matrix') print(confusion_matrix(y_test,y_pred))
```

```
print('accuracy') print(classification_report(y_test,y_pred))
```

OUTPUT:

```

print(y)
[5.1 3.7 1.5 0.4]
[4.6 3.6 1.  0.2]
[5.1 3.3 1.7 0.5]
[4.8 3.4 1.9 0.2]
[5.  3.  1.6 0.2]
[5.  3.4 1.6 0.4]
[5.2 3.5 1.5 0.2]
[5.2 3.4 1.4 0.2]
[4.7 3.2 1.6 0.2]
[4.8 3.1 1.6 0.2]
[5.4 3.4 1.5 0.4]
[5.2 4.1 1.5 0.1]
[5.5 4.2 1.4 0.2]
[4.9 3.1 1.5 0.2]
[5.  3.2 1.2 0.2]
[5.5 3.5 1.3 0.2]
[4.9 3.6 1.4 0.1]
[4.4 3.  1.3 0.2]
[5.1 3.4 1.5 0.2]
[5.  3.  1.  0.2]

In [21]: x_train, x_test, y_train, y_test = train_test_split(X,Y,test_size=0.3)

In [22]: classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(x_train, y_train)

Out[22]: KNeighborsClassifier
KNeighborsClassifier()

In [23]: y_pred=classifier.predict(x_test)

In [24]: print('confusion matrix')
print(confusion_matrix(y_test,y_pred))

confusion matrix
[[15  0  0]
 [ 0 17  2]
 [ 0  0 11]]

In [25]: print('accuracy')
print(classification_report(y_test,y_pred))

accuracy
precision    recall  f1-score   support

      0       1.00      1.00      1.00        15
      1       1.00      0.89      0.94        19
      2       0.85      1.00      0.92        11

   accuracy          0.96
  macro avg          0.95
weighted avg          0.96

```

Experiment 10

Implement the non-parametric Locally Weighted Regression algorithm in order to fit data points. Select appropriate data set for your experiment and draw graphs.

CODE:

from numpy import *

```
from os import listdir import
matplotlib import matplotlib.pyplot
as plt import pandas as pd import
numpy as np1 import numpy.linalg as
np from scipy.stats.stats import
pearsonr
```

```
def kernel(point,xmat, k): m,n =
np1.shape(xmat) weights =
np1.mat(np1.eye((m)))
for j in range(m):
diff = point - X[j]
weights[j,j] = np1.exp(diff*diff.T/(-2.0*k**2))
return weights
```

```
def localWeight(point,xmat,yamat,k):
wei = kernel(point,xmat,k)
W = (X.T*(wei*X)).I*(X.T*(wei*yamat.T))
return W
```

```
def localWeightRegression(xmat,yamat,k):
m,n = np1.shape(xmat)
ypred = np1.zeros(m)
for i in range(m): ypred[i] =
xmat[i]*localWeight(xmat[i],xmat,yamat,k) return
ypred
```

```
# load data points data =
pd.read_csv('tips.csv') bill =
```

```
np1.array(data.total_bill) tip =  
np1.array(data.tip)
```

```
#preparing and add 1 in bill mbill = np1.mat(bill) mtip = np1.mat(tip) # mat is used to  
convert to n dimesiona to 2 dimensional array form m= np1.shape(mbill)[1] # print(m)  
244 data is stored in m one = np1.mat(np1.ones(m))  
X= np1.hstack((one.T,mbill.T)) # create a stack of bill from ONE  
#print(X) #set k here ypred =  
localWeightRegression(X,mtip,2)  
SortIndex = X[:,1].argsort(0) xsort =  
X[SortIndex][:,0]
```

```
fig = plt.figure()  
ax = fig.add_subplot(1,1,1) ax.scatter(bill,tip, color='blue')  
ax.plot(xsort[:,1],ypred[SortIndex], color = 'red', linewidth=5)  
plt.xlabel('Total bill')  
plt.ylabel('Tip') plt.show()
```

```
import numpy as np from bokeh.plotting import figure,  
show, output_notebook from bokeh.layouts import  
gridplot from bokeh.io import push_notebook
```

```
def local_regression(x0, X, Y, tau):# add bias term x0 =  
np.r_[1, x0] # Add one to avoid the loss in information  
X = np.c_[np.ones(len(X)), X]  
# fit model: normal equations with kernel xw = X.T * radial_kernel(x0, X, tau) #  
XTranspose * W beta = np.linalg.pinv(xw @ X) @ xw @ Y #@ Matrix  
Multiplication or Dot Product
```

```
# predict value    return x0 @ beta # @ Matrix Multiplication or Dot
Product for prediction
```

```
def radial_kernel(x0, X, tau):
    return np.exp(np.sum((X - x0) ** 2, axis=1) / (-2 * tau * tau))
# Weight or Radial Kernel Bias Function
```

```
n = 1000
# generate dataset X = np.linspace(-3, 3, num=n) print("The
Data Set ( 10 Samples) X :\n",X[1:10]) Y = np.log(np.abs(X **
2 - 1) + .5) print("The Fitting Curve Data Set (10 Samples) Y
:\n",Y[1:10])
# jitter X
X += np.random.normal(scale=.1, size=n)
print("Normalised (10 Samples) X :\n",X[1:10]) domain =
np.linspace(-3, 3, num=300) print(" Xo Domain Space(10
Samples) :\n",domain[1:10])
```

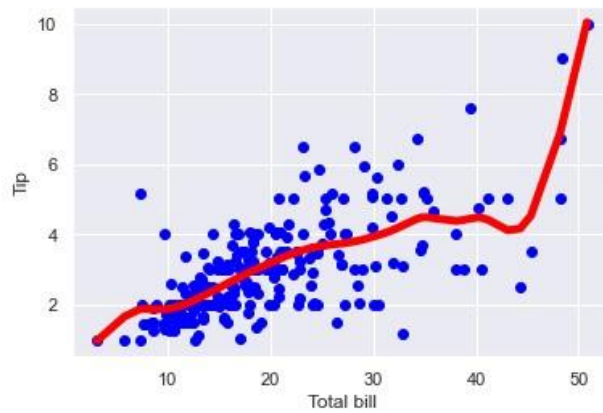
```
def plot_lwr(tau): # prediction through regression
prediction = [local_regression(x0, X, Y, tau) for x0 in domain]
plot = figure(plot_width=400, plot_height=400)
plot.title.text='tau=%g' % tau    plot.scatter(X, Y, alpha=.3)
plot.line(domain, prediction, line_width=2, color='red') return
plot
```

OUTPUT:

```

fig = plt.figure()
ax = fig.add_subplot(1,1,1)
ax.scatter(bill,tip, color='blue')
ax.plot(xsort[:,1],ypred[SortIndex], color = 'red', linewidth=5)
plt.xlabel('Total bill')
plt.ylabel('Tip')
plt.show()

```



The Data Set (10 Samples) X :

```

[-2.99399399 -2.98798799 -2.98198198 -2.97597598 -2.96996997 -2.96396396
 -2.95795796 -2.95195195 -2.94594595]

```

The Fitting Curve Data Set (10 Samples) Y :

```

[2.13582188 2.13156806 2.12730467 2.12303166 2.11874898 2.11445659
 2.11015444 2.10584249 2.10152068]

```

Normalised (10 Samples) X :

```

[-2.7984698 -3.00877009 -3.05888439 -2.95096415 -2.94588394 -2.97666794
 -3.01995 -3.08887995 -2.92471686]

```

Xo Domain Space(10 Samples) :

```

[-2.97993311 -2.95986622 -2.93979933 -2.91973244 -2.89966555 -2.87959866
 -2.85953177 -2.83946488 -2.81939799]

```

```

def plot_lwr(tau):

```